University of BRISTOL

Computer Science Year 2

Algorithms & Data

Estimation, Regression, Classification

Prof Alin Achim

University of BRISTOL

Last time …

- If a MVU estimator does not exist, or can not be found, the parameters can be obtained from the likelihood function.
- A Maximum Likelihood (ML) parameter estimate is found by maximising the likelihood function $p(x;\theta)$ which is essentially the probability of the data given the parameters;
- ML estimators are asymptotically efficient, as the number of observations increase and the covariance of the estimates tends to CRLB
- A major advantage of the MLE is that we can find an estimate from the given data numerically since it requires only the maximum of a known function. The Newton-Raphson iterative techniques or the Expectation-Maximisation (EM) algorithm can be used for iterative estimation of the parameters.

Asymptotic properties of the MLE

- As N → ∞ we have that $\hat{\theta} \to \theta$ (*consistent* estimator)

- Moreover:

$$E(\hat{\theta}) \to \theta$$
$$\text{var}(\hat{\theta}) \to CRLB$$

- The MLE is *asymptotically unbiased* and *asymptotically efficient*.

Example: Exponential distribution

- Assume $x_1, x_2, \ldots, x_N$ is a random sample from an exponential distribution:-

$$f(x|\theta) = \begin{cases} \theta e^{-\theta x}, & x \geq 0 \\ 0 & x < 0 \end{cases}$$

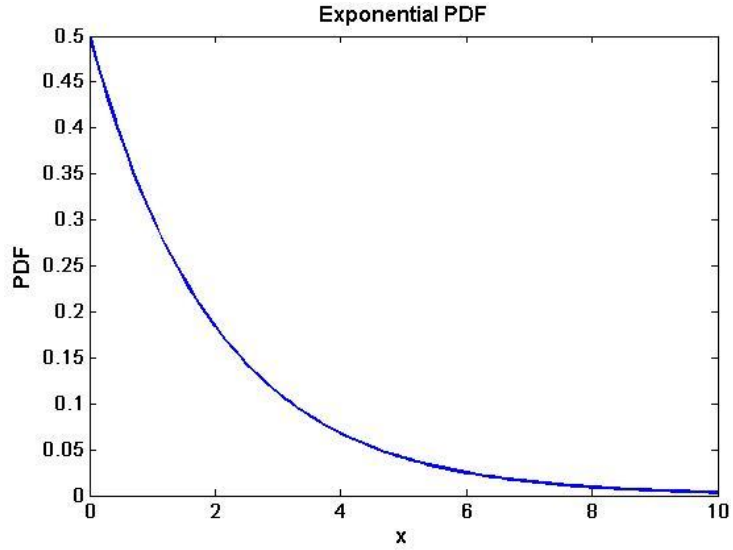

- The likelihood function is

$$L(\theta) = \prod_{i=1}^{N} f(x_i|\theta) = \theta^N e^{-\theta \sum_{i=1}^{N} x_i}$$

$$l(\theta) = N \log \theta - \theta \sum_{i=1}^{N} x_i$$

- Taking the derivative and setting equal to zero:-

$$\hat{\theta} = \frac{N}{\sum_{i=1}^{N} x_i} = \frac{1}{\bar{x}}$$

Joint MLE for several parameters

- Often in practice, a statistical model has more than one unknown parameter
- If there are k parameters, then we have a vector parameter $\boldsymbol{\theta}=(\theta_1, \theta_2, \ldots, \theta_k)^T$ and the PDF is written as $f(x|\boldsymbol{\theta})$.
- $\hat{\boldsymbol{\theta}} = \left(\hat{\theta}_1, \hat{\theta}_2, \ldots, \hat{\theta}_k\right)^T$ are the values in the parameter space that jointly maximise the likelihood function
- If $l(\boldsymbol{\theta})$ is differentiable then the vector estimate satisfy k joint differential equations

$$\frac{\partial}{\partial \theta_j} l(\theta_1, \theta_2, \ldots, \theta_k) = 0 \quad \text{for } j = 1,2,\ldots,k$$

Example: Normal distribution

- Assume $X \sim N(\mu, \sigma^2)$ i.e.

$$p(x|\mu,\sigma^2) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left[\frac{-(x-\mu)^2}{2\sigma^2}\right]$$
$$-\infty < \mu < \infty, \quad \sigma > 0$$

- The log-likelihood is

$$l(\mu,\sigma^2) = -\frac{N}{2}\log 2\pi - \frac{N}{2}\log\sigma^2 - \frac{\sum_{i=1}^{N}(x_i-\mu)^2}{2\sigma^2}$$

- Solving the likelihood equations yields:

$$\frac{\partial l}{\partial \mu} = -\frac{(-2)\sum_{i=1}^{N}(x_i-\mu)}{2\sigma^2} = 0 \Rightarrow \hat{\mu} = \bar{x}$$

$$\frac{\partial l}{\partial \sigma^2} = -\frac{N}{2\sigma^2} + \frac{1}{\sigma^4}\frac{\sum_{i=1}^{N}(x_i-\mu)^2}{2} = 0 \Rightarrow \hat{\sigma}^2 = \frac{\sum_i (x_i-\bar{x})^2}{N}$$

Objectives

- Least Square Estimation (LSE)
- Method of Moments (MoM)

Least Squares – The intuition

- Consider (again) the multiple observations:

$$x[n] = A + w[n], \text{ where } n = 0, 1, \ldots, N-1 \text{ and } w[n] \sim N(0, \sigma^2)$$

- Or more generally:

$x[n] = f(\theta_1, \theta_2, \ldots, \theta_M) + w[n]$, i.e. *f* is a function of M parameters

- The least square estimator (LSE) minimizes:

$$J = \sum_{n=0}^{N-1} \left(x[n] - f(\theta_1, \theta_2, \ldots, \theta_M) \right)^2$$

Least Squares - Properties

- LSE is widely used when estimating parameters for linear models
- No assumptions about the data are made
- If $w[n] \sim N(0, \sigma^2)$, LSE coincides with the MLE!
- Geometric interpretation: the LS estimate is an orthogonal projection of the data vector onto the space defined by the independent variable.

- Inverse problems formulation:
  - $\boldsymbol{b} = \boldsymbol{Ax} + \boldsymbol{n}, \boldsymbol{n} \sim N(0, \sigma^2)$
  - $\hat{\boldsymbol{x}} = min_{\boldsymbol{x}} \|\boldsymbol{b} - \boldsymbol{Ax}\|^2$

Geometric interpretation

- $\boldsymbol{A\hat{x}}$ is the orthogonal projection of $\boldsymbol{b}$ onto *range(A)*

$$\Leftrightarrow \boldsymbol{A}^T(\boldsymbol{b} - \boldsymbol{A\hat{x}}) = \boldsymbol{0} \Leftrightarrow \boldsymbol{A}^T\boldsymbol{A\hat{x}} = \boldsymbol{A}^T\boldsymbol{b}$$

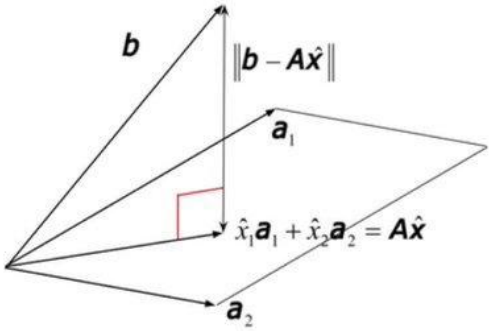

Least Squares – A much simpler example

- Suppose we have a random sample $\{x_n\}, n = 0, \ldots, N-1$, drawn from a population with mean μ_x and standard deviation σ_x.
  - We can express x_n using a linear model:

$$x_n = \mu_x + \varepsilon_n, E[\varepsilon_n] = 0 \; and \; E[\varepsilon_n^2] = \sigma_x^2$$

  - Estimate μ_x via LSE, that is minimise:

$$J(\mu_x) = \sum_{n=0}^{N-1} (x_n - \hat{\mu}_x)^2$$

$$\frac{\partial J(\mu_x)}{\partial \mu_x} = -2 \sum_{n=0}^{N-1} (x_n - \hat{\mu}_x) = 0$$

$$\Rightarrow \sum_{n=0}^{N-1} x_n = N\hat{\mu}_x \Rightarrow \hat{\mu}_x = \frac{1}{N} \sum_{n=0}^{N-1} x_n$$

Least Squares – A signal processing example

- One of the main applications of LSE is in Adaptive Noise Cancellation (ANC).

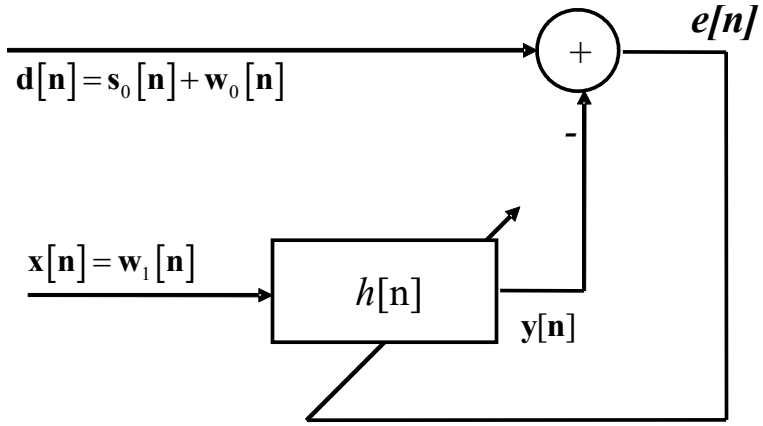

- A signal s_0 is received with some uncorrelated noise w_0. A second signal contains noise w_1 uncorrelated with s_0 but correlated with w_0 in some unknown way. The problem is to enhance the signal s_0 in the primary input.

- The principle of ANC is to estimate the noise in the primary input by filtering the reference noise signal with a linear filter. *How is the filter designed??*

Least Squares – A signal processing example

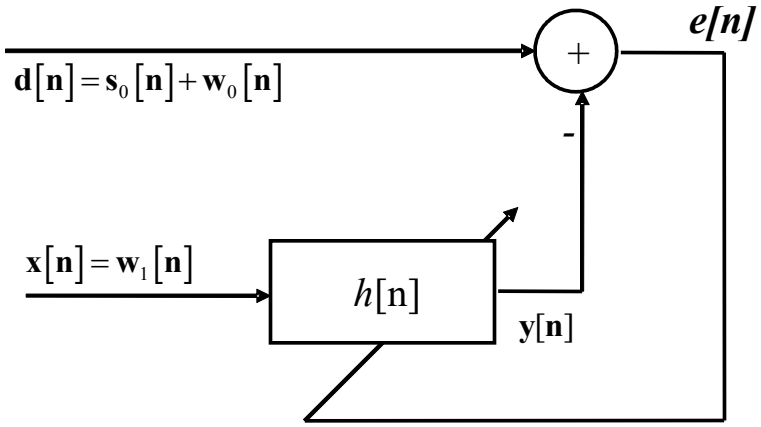

- Choose the weights at time *n* to minimize:

$$J[n] = \sum_{k=0}^{n} e^2[k] = \sum_{k=0}^{n} (d[k] - y[k])^2 = \sum_{k=0}^{n} \left(d[k] - \sum_{l=0}^{p-1} h[l]x[k-l] \right)^2$$

Method of Moments (MoM) – The intuition

- The method of moments seeks to equate the moments as implied by the underlying model of the population distribution $p(x)$ (mean, variance, skewness, kurtosis, etc) with the actual moments in the sample.

$$population\ moment = sample\ moment$$

- The population probability density function (pdf) is a function of the parameter we want to estimate, θ.

$$p(x) = f(x|\theta), \text{e.g.} f(x|\theta) \propto e^{-\frac{x^2}{2\theta^2}}$$

Method of Moments (MoM) – The intuition

$$population\ moment = sample\ moment$$

and

$$p(x) = f(x|\theta)$$

1. $E[x^n] = \int_{-\infty}^{\infty} x^n \, \mathrm{p}(x) = \int_{-\infty}^{\infty} x^n \, f(x|\theta) = g(\theta)$

2. $\theta = g^{-1}(E[x^n])$

3. $\theta = g^{-1}\left(n^{th}\ sample\ moment\ \right)$

MoM – Properties

- Main advantage: extremely easy to determine and implement
- The order of the moment(s) used depends on the parameter to be estimated, i.e. population mean/maximum/correlation is estimated using sample mean/maximum/correlation, etc
- To obtain estimates of several population parameters, several moments need to be employed
- As an example, consider again the case of the 2 parameter Gaussian distribution, in which case, we use the first and second order moments, i.e. the sample mean and sample variance:

$$E[x] = g_1(x|\theta_1) = \frac{1}{N}\sum_{n=1}^{N} x_n \; and \; E[x^2] = g_2(x|\theta_2) = \frac{1}{N}\sum_{n=1}^{N} x_n^2$$

MoM – Another simple example

- Suppose the population distribution follows an uniform distribution with unknown parameter θ:

$$p(X) = \begin{cases} \frac{1}{\theta}, for\ 0 \le X \le \theta \\ 0, \quad otherwise \end{cases}$$

- Given a random sample $\{x_n\}, n = 1, \dots, N$, estimate θ.

1. $E[X] = \int_0^\theta x \frac{1}{\theta} dx = \frac{\theta}{2}$
2. $\theta = 2E[X]$
3. $\hat{\theta} = 2\bar{X}$

$E[X^n] = g(\theta)$

$\theta = g^{-1}(E[X^n])$

$\hat{\theta} = g^{-1}\left(n^{th}\ sample\ moment\right)$

Real world application – image fusion

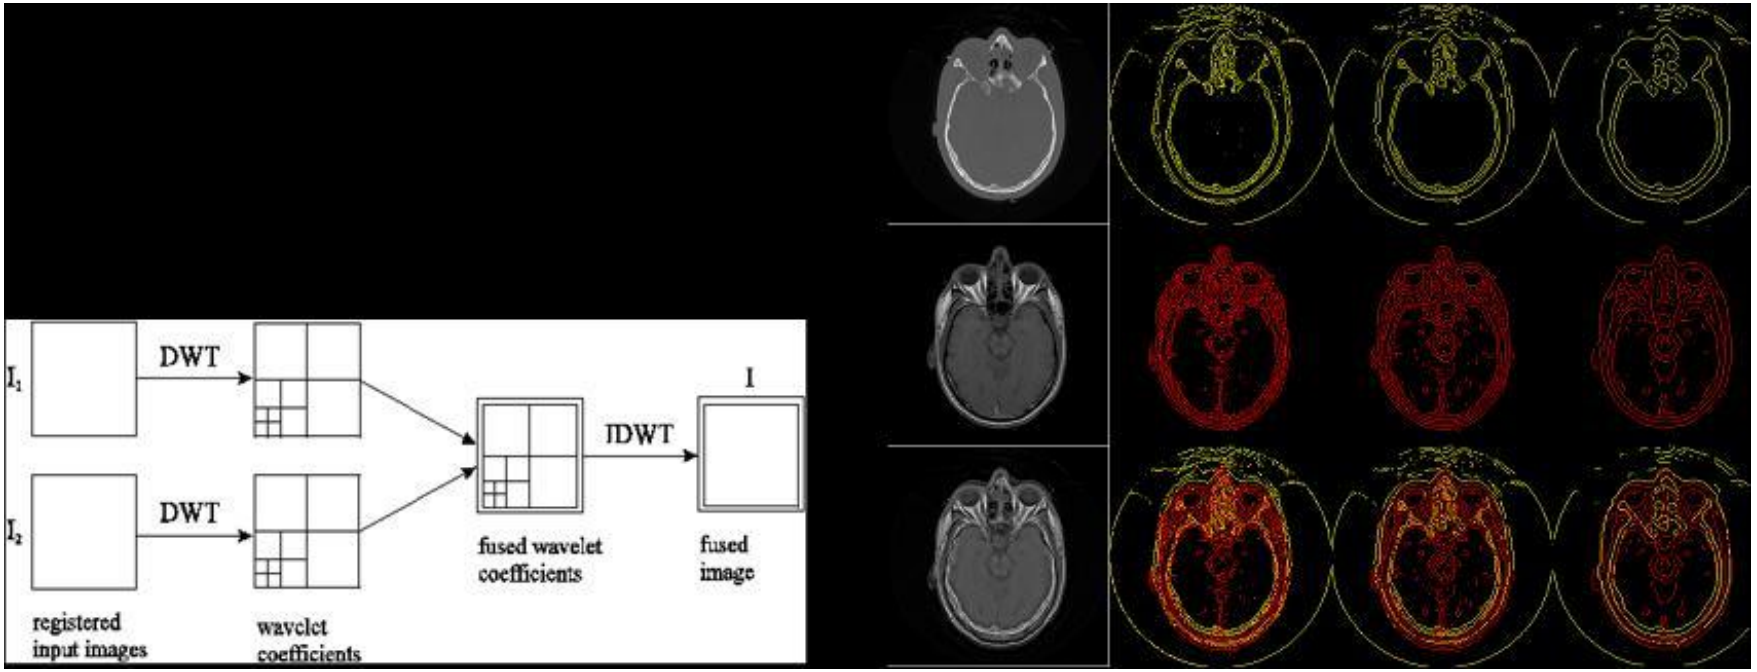

Image fusion algorithms

- Pixel – level fusion
- Region – level fusion
- Compressive fusion

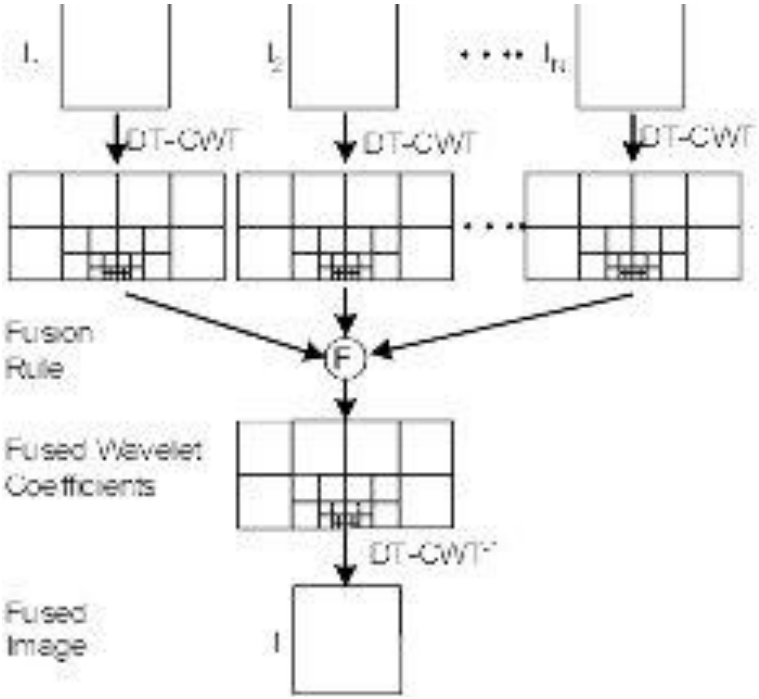

Pixel level approaches:

- Maximum selection scheme
- Average scheme
- Weighted average scheme

The FLOM weighted average (FLOM-WA) scheme

I. **DTCWT of both images**

II. **For each subband pair**

1. **Estimate γ_{x_1} and γ_{x_2}**
2. **Compute $\mathrm{Corr}_{\alpha}(X_1, X_2)$**
3. **Calculate the fused coefficient using**

$$Y = w_1 X_1 + w_2 X_2$$

III. **Average coefficients in lowpass residual**

IV. **IDTCWT**

Example WA-FLOM Fusion of MR & CT images

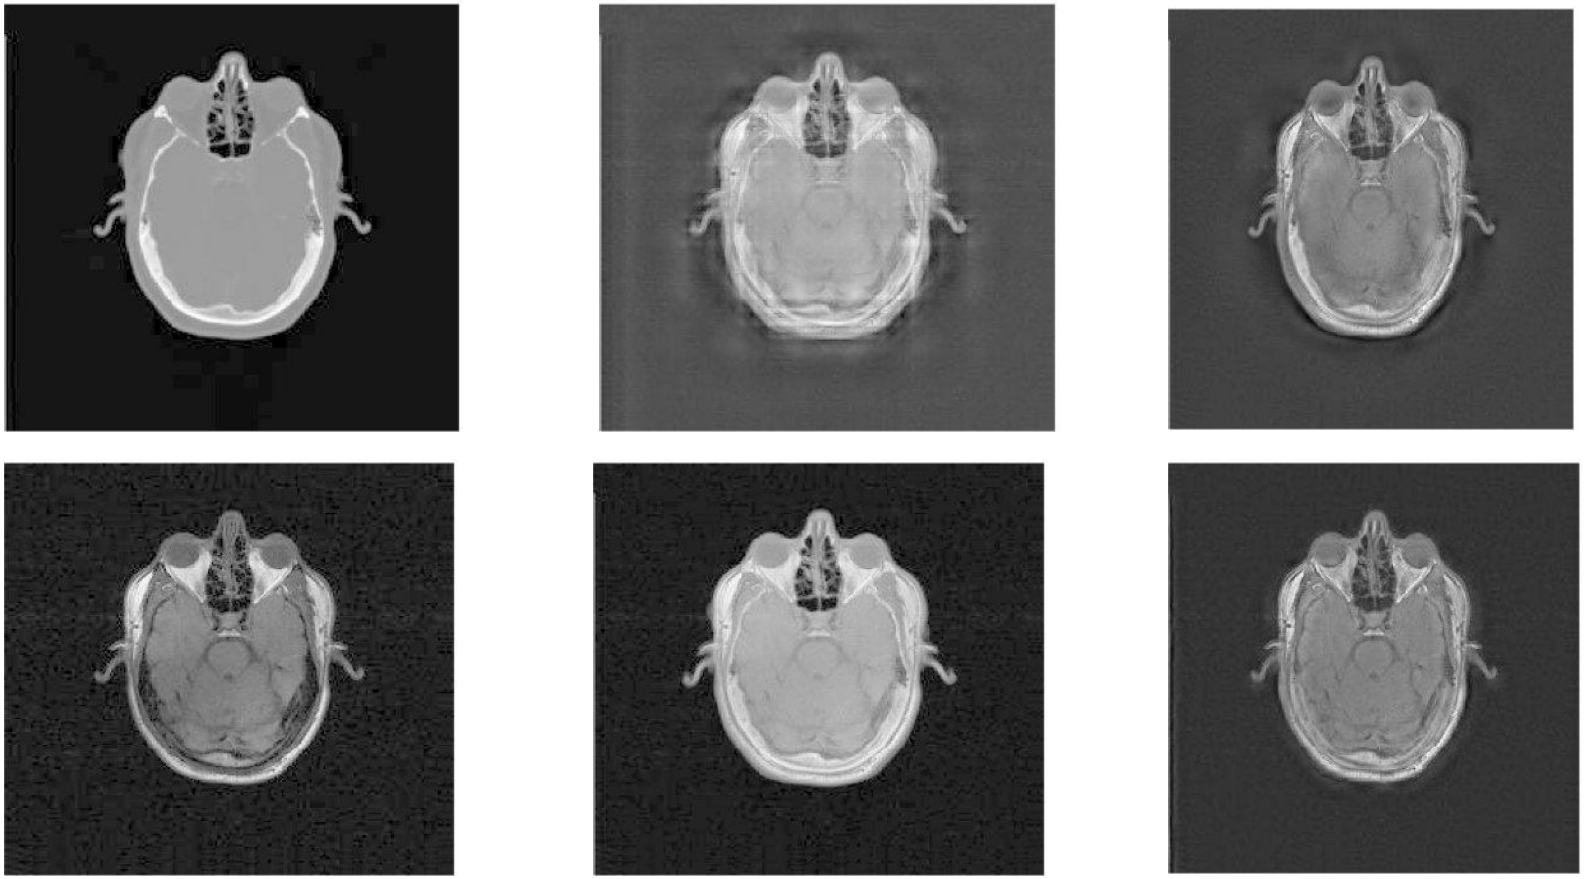

Summary of LSE & MoM

- LSE:
  - minimises the sum of squares between the measurements and the model
  - no assumption about the data — generally applicable estimators
  - best in the context of linear models
- MoM:
  - equate the sample with population moments
  - the simplest, intuitive, works well in straightforward cases
  - estimators not always with good properties (e.g. for small sample size)